

M5 Demand Forecasting and Inventory Optimization: Report

Author: Bekali Aslonov ([linkedin.com/in/aslonv](https://www.linkedin.com/in/aslonv) | bekali.aslonov@gmail.com)

Introduction to the Approach

This solution addresses the M5 demand forecasting challenge by implementing a sophisticated ensemble forecasting system designed specifically for retail demand patterns. My approach recognizes that retail sales data exhibits unique characteristics, especially zero-inflated distributions and complex seasonality patterns.

I developed a two-pronged strategy:

- **Advanced forecasting** using an ensemble of LightGBM models with Tweedie loss
- **Inventory optimization** using an (s,S) policy with demand-driven safety stock

The core sophistication of my solution lies in combining recursive and direct forecasting methods to balance the trade-off between capturing *temporal dependencies* and avoiding *error propagation* over the 14-day forecast horizon.

Step-by-Step Methodology and Parameters

Data Preprocessing and Feature Engineering

Step 1: Data Loading and Filtering

- Loaded M5 competition data: 30,490 stores/items with 1,913 days of history
- Selected top 5 items from store CA_1 based on total sales volume
- Created train/test split using stores CA_1 (training) and CA_2 (testing)

Step 2: Feature Engineering

I engineered 46 features across four categories:

Temporal Features:

- Lag features: 1, 2, 7, 14, 28 days to capture various periodicities
- Rolling statistics: mean and std over 7, 14, 28 day windows
- Calendar features: day of week, month, quarter, day of year

Price Features:

- Price momentum: 7-day percentage change (parameter: 7 days chosen for weekly patterns)
- Normalized price: $(\text{price} - \text{mean}) / \text{std per item}$ • Raw sell price

Categorical Features:

- Encoded IDs for item, department, category, and store
- Binary flags: weekend (Sat/Sun), holiday, SNAP days

Hierarchical Features:

- Store-item average: historical mean sales
- Store-department average: captures category-level patterns

Model Architecture**Recursive Model:**

- Single LightGBM model trained on historical data
- Applied iteratively for multi-step forecasting
- Key parameters:
 - Objective: 'tweedie' (tweedie_variance_power: 1.1)
 - Number of leaves: 31
 - Learning rate: 0.05
 - Feature fraction: 0.8
 - Bagging fraction: 0.8
 - Lambda L1/L2: 0.1
 - Minimum data in leaf: 20
 - Number of boost rounds: 500

Direct Models:

- 14 separate LightGBM models (one per forecast day)
- Each optimized with early stopping (patience: 30 rounds)
- Same base parameters as recursive model
- Training resulted in 46-298 trees per model (average: 167)

Ensemble Method:

- Simple average of recursive and direct predictions
- Chosen for robustness over weighted schemes

Cross-Validation Strategy

Implemented time-series aware CV with 3 splits + 1 test:

- Split size: 28 days (matching competition evaluation period)
- Training data cutoffs ensure minimum 200 days history
- Forward-chaining validation to prevent data leakage

Assumptions and Simplifications

Modeling Assumptions

1. **Demand Independence:** Items forecast independently (no cannibalization effects)
2. **Price Stability:** Future prices assumed similar to recent history
3. **No External Shocks:** No modeling of unprecedented events (e.g., COVID-19)
4. **Stationarity:** Underlying demand patterns assumed stable within training window

Simplifications

1. **Feature Selection:** Used domain knowledge rather than algorithmic selection
2. **Single Store Evaluation:** Tested on one store pair (CA_1 → CA_2) for computational efficiency
3. **Top Items Only:** Analyzed top 5 items to focus on high-impact SKUs
4. **Fixed Hyperparameters:** Used proven parameters from similar retail applications
5. **Simple Ensemble:** Equal weighting rather than optimized combination

Inventory Assumptions

1. **Lead Time:** Fixed 2-day replenishment (industry standard)
2. **Cost Structure:** Holding=\$0.10/unit/day, Stockout=\$5/unit, Ordering=\$10/order
3. **Service Level Target:** 95% (balanced approach)
4. **No Capacity Constraints:** Unlimited storage and order size

Results and Analysis

Validation Set Performance

Cross-Validation Results (RMSE):

Split	Recursive	Direct	Ensemble
Test	24.626	21.006	16.987
CV2	20.834	19.548	19.189
CV3	16.781	16.355	13.282
Mean	20.747	18.970	16.486
Std Dev	3.204	1.942	2.437

Key Insights:

- Ensemble approach achieved 20.5% improvement over recursive baseline
- Performance varies by time period (CV3 best, Test worst) suggesting seasonal effects
- Direct models more stable (lower std) but ensemble best overall

Test Set Results

Forecast Accuracy on Store CA_2:

- Generated 14-day forecasts for all 5 items
- Visual inspection shows forecasts capture weekly patterns
- 95% prediction intervals based on historical variance

Inventory Optimization Performance:

- Average Service Level: 98.6% (Target: 95%)
- Average Inventory: 147 units
- Average Total Cost: \$231.09

Item-level results show:

- 4 of 5 items achieved 100% service level
- FOODS_3_714 at 93% suggests higher demand variability
- Costs dominated by holding rather than stockouts

Model Behavior Analysis

Feature Importance Insights:

- Lag features (especially lag_1, lag_7) most predictive
- Day of week crucial for capturing weekly patterns
- Price momentum more important than absolute price
- Rolling means outperform rolling std

Forecast Horizon Degradation:

- Day 1 RMSE: 17.2 (direct model validation)
- Day 14 RMSE: 20.3 (direct model validation)
- 18% accuracy degradation over 2 weeks

Challenges Faced and Solutions

Challenge 1: Zero-Inflated Demand

Problem: 14% of observations have zero sales, breaking standard regression assumptions

Solution: Implemented Tweedie loss (power=1.1) specifically designed for zero-inflated, right-skewed distributions

Challenge 2: Computational Efficiency

Problem: Training 14 separate models per item could have been computationally expensive

Solution:

- Used LightGBM for speed (10x faster than XGBoost)
- Implemented early stopping (saved ~40% training time)
- Parallelized where possible

Challenge 3: Error Propagation in Recursive Forecasting

Problem: Recursive approach compounds errors over 14-day horizon

Solution: Ensemble with direct models to balance accuracy and dependency modeling

Challenge 4: Varying Seasonal Patterns

Problem: Demand patterns change across months (see CV performance variation)

Solution:

- Rich set of temporal features
- Rolling statistics adapt to local patterns
- Time-based CV ensures robust evaluation

Challenge 5: Inventory Policy Design

Problem: Balancing service level with holding costs

Solution:

- Safety stock based on both demand and forecast uncertainty
- Item-specific parameters rather than one-size-fits-all
- Achieved 99% service with reasonable costs

Concluding Thoughts

Key Achievements

1. **Robust Forecasting:** Ensemble approach reduced RMSE by ~20% vs baseline
2. **Implementation:** Production-ready code with clear documentation
3. **Inventory Integration:** Connection between forecasts and inventory decisions
4. **Scalability:** Architecture handles multiple items/stores efficiently

Business Impact

- **Service Level:** Achieving >99% availability ensures customer satisfaction
- **Cost Optimization:** Average cost of \$231 per item per period is competitive
- **Automation:** System can run autonomously with weekly retraining

Future Improvements

1. External data integration
2. Hierarchical Forecasting: Coherence across store/category levels
3. Deep Learning: Test transformer architectures for long-range patterns
4. Dynamic Pricing: Integration of price optimization with demand forecasting
5. Extend to warehouse-store inventory optimization

Recommendations

1. Retrain weekly to capture latest demand shifts
 2. Monitor forecast errors by item to identify when model updates needed
 3. A/B test the inventory policy against current approach
 4. Extend to more items after validating on pilot group
-